

Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

3 **Anonymous author**

4 Anonymous affiliation

5 **Anonymous author**

6 Anonymous affiliation

7 — Abstract —

8 We benchmark four proof assistants (Agda, Idris 2, Lean 4 and Rocq) through a single test suite.
9 We focus our benchmarks on the basic features that all systems based on a similar foundations
10 (dependent type theory) have in common. We do this by creating an “over language” in which to
11 express all the information we need to be able to output *correct and idiomatic syntax* for each of our
12 targets. Our benchmarks further focus on “basic engineering” of these systems: how do they handle
13 long identifiers, long lines, large records, large data declarations, and so on.

14 Our benchmarks reveals both flaws and successes in all systems. We give a thorough analysis of
15 the results.

16 We also detail the design of our extensible system. It is designed so that additional tests and
17 additional system versions can easily be added. A side effect of this work is a better understanding
18 of the common abstract syntactic structures of all four systems.

19 **2012 ACM Subject Classification** Replace ccsdesc macro with valid one

20 **Keywords and phrases** Add keywords

21 **Digital Object Identifier** 10.4230/LIPIcs.CVIT.2016.23

22 **1 Introduction**

23 Production-grade implementations of dependently typed programming languages are complica-
24 ted pieces of software that implement many intricate and potentially expensive algorithms.
25 As such, large amounts of engineering effort has been dedicated to optimizing these com-
26 ponents. Unfortunately, engineering time is a finite resource, and this necessarily means
27 that other parts of these systems get comparatively less attention. This often results in
28 easy-to-miss performance problems: we have heard anecdotes from a proof assistant developer
29 that a naïve $O(n^2)$ fresh name generation algorithm used for pretty-printing resulted in 100x
30 slowdowns in some pathological cases.

31 This suggests that a benchmarking suite that focuses on these simpler components
32 could reveal some (comparatively) easy potential performance gains. Moreover, such a
33 benchmarking suite would also be valuable for developers of new dependently typed languages,
34 as it is much easier to optimize with a performance goal in mind. This is an instance of
35 the classic $m \times n$ language tooling problem: constructing a suite of m benchmarks for n
36 languages directly requires a quadratic amount of work up front, and adding either a new
37 test case or a new language to the suite requires an additional linear amount of effort.

38 Like most $m \times n$ tooling problems, the solution is to introduce a mediating tool. In our
39 case, we ought to write all of the benchmarks in an intermediate language, and then translate
40 that intermediate language to the target languages in question. There are existing languages
41 like Dedukti[2] or Informath[1] that attempt to act as an intermediary between popular proof
42 assistants, but these tools typically focus on translating the *content* of proofs, not exact
43 syntactic structure. To fill this gap, we have created the Panbench system, which consists of:



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:10



Leibniz International Proceedings in Informatics

LIPIcs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

23:2 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

- 44 1. An extensible embedded Haskell DSL that encodes the concrete syntax of a typical
45 dependently typed language.
- 46 2. A series of compilers for that DSL to Agda, Idris 2, Lean 4, and Rocq.
- 47 3. A benchmarking harness that can perform sandboxed builds of multiple revisions Agda,
48 Idris 2, Lean 4, Rocq.
- 49 4. An incremental build system that can produce benchmarking reports as static HTML
50 files or PGF plots¹.

51 2 Methodology

52 This is really documenting the 'experiment'. The actual details of the thinking behind
the design is in Section 6.

53 this itemized
list should be
expanded into
actual text

- 54 └─ single language of tests
55 └─ document the setup of tests, high level
56 └─ document the setup of testing infrastructure, high level
57 └─ linear / exponential scaling up of test 'size'

58 3 Results

59 4 Discussion

60 5 Infrastructure

61 One of the major goals of Panbench is to make performance analysis as low-cost as possible for
62 language developers. Meeting this goal requires a large amount of supporting infrastructure:
63 simply generating benchmarks is not very useful if you cannot run them nor analyze their
64 outcomes. After some discussion, we concluded that any successful language benchmarking
65 system should meet the following criteria:

- 66 1. It must provide infrastructure for performing sandboxed builds of compilers from source.
67 Asking potential users to set up four different toolchains presents an extremely large
68 barrier to adoption. Moreover, if we rely on user-provided binaries, then we have no hope
69 of obtaining reproducible results, which in turn makes any insights far less actionable.
- 70 2. It must allow for multiple revisions of the same tool to be installed simultaneously. This
71 enables developers to easily look for performance regressions, and quantify the impact of
72 optimizations.
- 73 3. It must allow for multiple copies of the *same* version tool to be installed with different
74 build configurations. This allows developers to look for performance regressions induced
75 by different compiler versions/optimizations.
- 76 4. It must be able to be run locally on a developer's machine. Cloud-based tools are often
77 cumbersome to use and debug, which in turn lowers adoption.
- 78 5. It must present a declarative interface for creating benchmarking environments and
79 running benchmarks. Sandboxed builds of tools are somewhat moot if we cannot trust
80 that a benchmark was run with the correct configuration.

¹ All plots in this paper were produced directly by Panbench!

81 6. It must present performance results in a self-contained format that is easy to understand
82 and share. Performance statistics that require large amounts of post-processing or
83 dedicated tools to view can not be easily shared with developers, which in turn makes
84 the data less actionable.

85 Of these criteria, the first four present the largest engineering challenge, and are tantamount to constructing a meta-build system that is able to orchestrate *other* build systems.
86 We approached the problem by constructing a bespoke content-addressed system atop of
87 Shake [8], which we discuss in section 5.1. The final two criteria also presented some unforeseen
88 difficulties, which we detail in 5.2.

90 5.1 The Panbench Build System

91 As noted earlier, we strongly believe that any benchmarking system should provide infrastructure
92 for installing multiple versions of reproducibly built software. Initially, we intended
93 to build this infrastructure for Panbench atop of Nix [4]. This is seemingly a perfect fit;
94 after all, Nix was designed to facilitate almost exactly this use-case. However, after further
95 deliberation, we came to the conclusion that Nix did not quite meet our needs for the
96 following reasons:

- 97 1. Nix does not work natively on Windows. Performance problems can be operating system
98 specific, so ruling out an OS that has a large user base that is often overlooked in testing
99 seems unwise².
- 100 2. Nix adds a barrier to adoption. Installing Nix is a somewhat invasive process, especially
101 on MacOS³. We believe that it is somewhat unreasonable to ask developers to add users
102 and modify their root directory to run a benchmarking tool, and strongly suspect that
103 this would hamper adoption.

104 With the obvious option exhausted, we opted to create our own Nix-inspired build system
105 based atop Shake [8]. This avoids the aforementioned problems with Nix: Shake works on
106 Windows, and only requires potential users to install a Haskell toolchain.

107 The details of content-addressed build systems are a deep topic unto themselves, so we
108 will only describe the key points. Systems like Nix use an *input-addressing* scheme, wherein
109 the results of a build are stored on disk prefixed by a hash of all build inputs. Crucially,
110 this lets the build system know where to store the result of the build before the build is run,
111 which avoids vicious cycles where the result of a build depends on its own hash. However,
112 most input-addressed systems also require that the hash of the inputs *solely* determines the
113 output. On its face, this is a reasonable ask, but taking this idea to its logical conclusion
114 requires one to remove *any* dependency on the outer environment, which in turn forces one
115 to re-package the entirety of the software stack all the way down to `libc`. This is an admirable
116 goal in its own right, but is actually somewhat counterproductive for our use case: there
117 is a very real chance that we might end up benchmarking our sub-par repackaging of some
118 obscure C dependency four layers down the stack.

119 To avoid this cascading series of dependency issues, Panbench takes a hybrid approach,
120 wherein builds are input-addressed, but are allowed to also depend on the external environment.
121 Additionally, the results of builds are also hashed, and stored out-of-band inside of a

² Currently, Panbench does not support Windows, but this is an artifact of prioritization, and not a fundamental restriction.

³ The situation is even worse on x86-64 Macs, which most Nix installers simply do not support.

23:4 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

122 Shake database. This hash is used to perform invalidate downstream results, and also as a
123 fingerprint to identify if two benchmarks were created from the same binary. This enables
124 a pay-as-you go approach to reproducibility, which we hope will result in a lower adoption
125 cost. Moreover, we can achieve fully reproducible builds by using Nix as a meta-meta build
126 system to compile Panbench: this is how we obtained the results presented in Section 3.

127 5.2 Running Benchmarks and Generating Reports

128 As noted in the introduction to this section, we believe that benchmarking tools should
129 present a *declarative* interface for writing not just single benchmarking cases, but entire
130 benchmarking suites and their corresponding environments. Panbench accomplishes this
131 by *also* implementing the benchmark execution framework atop Shake. This lets us easily
132 integrate the process of tool installation with environment setup, but introduces its own set
133 of engineering challenges.

134 The crux of the problem is that performance tests are extremely sensitive to the current
135 load on the system. This is largely at odds with the goals of a build system, which is to
136 completely saturate all system resources to try to complete a build as fast as possible. This
137 can be avoided via careful use of locks, but we are then faced with another, larger problem.
138 Haskell is a garbage collected language, and running the GC can put a pretty heavy load on
139 the system. Moreover, the GHC runtime system is very well engineered, and is free to run
140 the garbage collector inside of `safe` FFI calls, and waiting for process completion is marked
141 as `safe`.

142 To work around this, we take opt to eschew existing process interaction libraries, and
143 implement the benchmark spawning code in C⁴. This lets us take the rather extreme step of
144 linking against the GHC runtime system so that we can call `rts_pause`, which pauses all
145 other Haskell threads and GC sweeps until `rts_resume` is called.

146 Initially, we thought that this was the only concern that would arise by tightly integrating
147 the build system with the benchmark executor. However, our initial benchmarks on Linux
148 systems displayed some very strange behaviour, wherein the max resident set size reported
149 by `getrusage` and `wait4` would consistently always report a reading of approximately 2
150 gigabytes halfway through a full benchmarking suite. After some investigating, we discovered
151 the Linux preserves resource usage statistics across calls to `execve`. Consequentially, this
152 means that we are unable to measure any max RSS that is lower than max RSS usage of
153 Panbench itself. Luckily, our lowest baseline is Agda at 64 megs, and we managed to get the
154 memory usage of Panbench itself down to 10 megs via some careful optimization and GC
155 tuning.

156 Currently, the statistics that Panbench gathers can then be rendered into standalone
157 HTML files with `vega-lite` plots, or into `TeX`files containing PGF plots. We intend to add
158 more visualization and statistic analysis tools as the need arises.

159 6 The Design of Panbench

160 At its core, Panbench is a tool for producing grammatically well-formed concrete syntax
161 across multiple different languages. Crucially, Panbench does *not* require that the syntax
162 produced is well-typed or even well-scoped: if it did, then it would be impossible to benchmark
163 how systems perform when they encounter errors. This seemingly places Panbench in stark

⁴ This is why Panbench does not currently support Windows.

¹⁶⁴ contrast with other software tools for working with the meta-theoretic properties of type
¹⁶⁵ systems, which are typically concerned only with well-typed terms.

¹⁶⁶ However, core task of Panbench is not that different from the task of a logical framework [5]:
¹⁶⁷ Both systems exist to manipulate judgements, inference rules, and derivations: Panbench
¹⁶⁸ just works with *grammatical* judgements and production rules rather than typing judgments
¹⁶⁹ and inference rules. In this sense Panbench is a *grammatical* framework⁵ rather than a logical
¹⁷⁰ one.

¹⁷¹ This similarity let us build Panbench atop well-understood design principles. In particular,
¹⁷² a mechanized logical framework typically consists of two layers:

- ¹⁷³ 1. A layer for defining judgements à la relations.
- ¹⁷⁴ 2. A logic programming layer for synthesizing derivations.

¹⁷⁵ To use a logical framework, one first encodes a language by laying out all of the judgements.
¹⁷⁶ Then, one needs to prove an adequacy theorem on the side that shows that their encoding of
¹⁷⁷ the judgements actually aligns with the language. However, if one wanted to mechanize this
¹⁷⁸ adequacy proof, then a staged third layer that consists of a more traditional proof assistant
¹⁷⁹ would be required.

¹⁸⁰ If we take this skeleton of a design and transpose it to work with grammatical constructs
¹⁸¹ rather than logical ones, we will also obtain three layers:

- ¹⁸² 1. A layer for defining grammars à la relations.
- ¹⁸³ 2. A logic programming layer for synthesizing derivations.
- ¹⁸⁴ 3. A staged functional programming layer for proving “adequacy” results.

¹⁸⁵ In this case, an adequacy result for a given language \mathcal{L} is a constructive proof that all
¹⁸⁶ grammatical derivations written within the framework can be expressed within the concrete
¹⁸⁷ syntax of a language \mathcal{L} . However, the computational content of such a proof essentially
¹⁸⁸ amounts to a compiler written in the functional programming layer.

¹⁸⁹ 6.1 Implementing The Grammatical Framework

¹⁹⁰ Implementing a bespoke hybrid of a logic and functional programming language is no small
¹⁹¹ feat, and also requires prospective users to learn yet another single-purpose tool. Luckily,
¹⁹² there already exists a popular, industrial-grade hybrid logic/functional programming language
¹⁹³ in wide use: GHC Haskell.

¹⁹⁴ At first glance, Haskell does not contain a logic programming language. However, if we
¹⁹⁵ enable enough GHC extensions, the typeclass system can be made *just* rich enough to encode
¹⁹⁶ a simply-typed logical framework. The key insight is that we can encode each production
¹⁹⁷ rule using multi-parameter type classes with a single method. Moreover, we can encode our
¹⁹⁸ constructive adequacy proofs for a given set of production rules as instances that translate
¹⁹⁹ each of the productions in the abstract grammar to productions in the syntax of an actual
²⁰⁰ language.

²⁰¹ As a concrete example, consider the grammar of the following simple imperative language.

²⁰² $\langle \text{expr} \rangle := \text{x}$
²⁰³ | n

Reed: Cite something

⁵ Not to be confused with *the* Grammatical Framework [9], which aims to be a more natural-language focused meta-framework for implementing and translating between grammars.

23:6 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

```

205   |  ⟨expr⟩ ‘+’ ⟨expr⟩
206   |  ⟨expr⟩ ‘*’ ⟨expr⟩

208 ⟨stmt⟩ := ⟨var⟩ ‘=’ ⟨expr⟩
209   |  ‘while’ ⟨expr⟩ ‘do’ ⟨stmt⟩
210   |  ⟨stmt⟩ ‘;’ ⟨stmt⟩

```

211 We can then encode this grammar with the following set of multi-parameter typeclasses:

Listing 1 An example tagless encoding.

```

212 class Var expr where
213   var :: String → expr
214
215 class Lit expr where
216   lit :: Int → expr
217
218 class Add expr where
219   add :: expr → expr → expr
220
221 class Mul expr where
222   mul :: expr → expr → expr
223
224 class Assign expr stmt where
225   assign :: String → expr → stmt
226
227 class While expr stmt where
228   while :: expr → stmt → stmt
229
230 class AndThen stmt where
231   andThen :: stmt → stmt → stmt
232
233

```

234 This style of language encoding is typically known as the untyped variant of *finally*
 235 *tagless*[3], and is well-known technique. However, our encoding is a slight refinement of
 236 the usual tagless style. In particular, we restrict ourselves to a single class per production
 237 rule, whereas other tagless encodings often use a class per syntactic category. This more
 238 fine-grained approach allows us to encode grammatical constructs that are only supported
 239 by a subset of our target grammars; see section 6.2 for examples.

240 Unfortunately, the encoding above has some serious ergonomic issues. In particular,
 241 expressions like `assign "x"` (`lit 4`) will result in an unsolved metavariable for `expr`, as there
 242 may be multiple choices of `expr` to use when solving the `Assign ?expr stmt` constraint. Luckily,
 243 we can resolve ambiguities of this form through judicious use of functional dependencies[6],
 244 as demonstrated below.

Listing 2 A tagless encoding with functional dependencies.

```

245 class Assign expr stmt | stmt → expr where
246   assign :: String → expr → stmt
247
248 class While expr stmt | stmt → expr where
249   while :: expr → stmt → stmt
250

```

252 **6.2 Designing The Language**

253 Now that we've fleshed out how we are going to encode our grammatical framework into
254 our host language, it's time to design our idealized abstract grammar. All of our target
255 languages roughly agree on a subset of the grammar of non-binding terms: the main sources
256 of divergence are binding forms and top-level definitions⁶. This is ultimately unsurprisingly:
257 dependent type theories are fundamentally theories of binding and substitution, so we would
258 expect some variation in how our target languages present the core of their underlying
259 theories.

260 This presents an interesting language design problem. Our idealized grammar will need to
261 find some syntactic overlap between all of our four target languages. Additionally, we would
262 also like for our solution to be (reasonably) extensible. Finding the core set of grammatical
263 primitives to accomplish this task is surprisingly tricky, and requires a close analysis of fine
264 structure of binding.

265 **6.2.1 Binding Forms**

266 As users of dependently typed languages are well aware, a binding form carries much more
267 information than just a variable name and a type. Moreover, this extra information can have
268 a large impact on typechecking performance, as is the case with implicit/visible arguments.
269 To make matters worse, languages often offer multiple syntactic options for writing the same
270 binding form, as is evidenced by the prevalence of multi-binders like $(x\ y\ z : A) \rightarrow B$. Though
271 such binding forms are often equivalent to their single-binder counterparts as *abstract* syntax,
272 they may have different performance characteristics, so we cannot simply lower them to a
273 uniform single-binding representation. To account for these variations, we have designed a
274 sub-language dedicated solely to binding forms. This language classifies the various binding
275 features along three separate axes: binding arity, binding annotations, and binding modifiers.

276 Binding arities and annotations are relatively self-explanatory, and classify the number of
277 names bound, along with the type of annotation allowed. Our target languages all have their
278 binding arities falling into one of three classes: *n*-ary, unary, or nullary. We can similarly
279 characterise annotations into three categories: required, optional, or forbidden.

280 This language of binders is encoded in the implementation as a single class `Binder` that is
281 parameterised by higher-kinded types `arity :: Type -> Type` and `ann :: Type -> Type`, and
282 provide standardized types for all three binding arities and annotations.

Listing 3 The basic binding constructs in Panbench.

```
283 class Binder arity nm ann tm cell | cell → nm tm where
284   binder :: arity nm → ann tm → cell
285
286   -- | No annotation or arity.
287   data None nm = None
288
289   -- | A single annotation or singular arity.
290   newtype Single a = Single { unSingle :: a }
291
292   -- | Multi-binders.
293   type Multi = []
```

⁶ As we shall see in section 6.2.2, the syntactical divergence of top-level definitions is essentially about binders as well.

23:8 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

```

295   -- | Infix operator for an annotated binder with a single name.
296   (.:) :: (Binder Single nm Single tm cell) => nm → tm → cell
297   nm .: tp = binder (Single nm) (Single tp)
298
299   -- | Infix operator for an annotated binder.
300   (.:*) :: (Binder arity nm Single tm cell) => arity nm → tm → cell
301   nms .: tp = binder nms (Single tp)
302

```

304 Production rules that involve binding forms are encoded as classes that are parametric over
 305 a notion of a binding cell, as demonstrated below.

Listing 4 The Panbench class for II-types.

```

306 class Pi cell tm | tm → cell where
307   pi :: [cell] → tm → tm
308

```

310 Decoupling the grammar of binding forms from the grammar of binders themselves allows
 311 us to be somewhat polymorphic over the language of binding forms when writing generators.
 312 This in turn means that we can potentially re-use generators when extending Panbench with
 313 new target grammars that may support only a subset of the binding features present in our
 314 four target grammars.

315 Binding modifiers, on the other hand, require a bit more explanation. A binding modifier
 316 captures features like implicit arguments, which do not change the number of names bound
 317 nor their annotations, but rather how those bound names get treated by the rest of the
 318 system. Currently, Panbench only supports visibility-related modifiers, but we have designed
 319 the system so that it is easy to extend with new modifiers; EG: quantities in Idris 2 or
 320 irrelevance annotations in Agda.

321 The language of binding modifiers is implemented as the following set of Haskell type-
 322 classes.

Listing 5 Typeclasses for binding modifiers.

```

323 class Implicit cell where
324   implicit :: cell → cell
325
326 class SemiImplicit cell where
327   semiImplicit :: cell → cell
328

```

330 This is a case where the granular tagless encoding shines. Both Lean 4 and Rocq have a
 331 form of semi-implicits⁷, whereas Idris 2 and Agda have no such notion. Decomposing the
 332 language of binding modifiers into granular pieces lets us write benchmarks that explicitly
 333 require support for features like semi-implicits. Had we used a monolithic class that encodes
 334 the entire language of modifiers, we would have to resort to runtime errors (or, even worse,
 335 dubious attempts at translation).

336 6.2.2 Top-Level Definitions

337 The question of top-level definitions is much thornier, and there seems to be less agreement
 338 on how they ought to be structured. Luckily, we can re-apply many of the lessons we

⁷ We use the term *semi-implicit* argument to refer to an implicit argument that is not eagerly instantiated. In Rocq these are known as non-maximal implicits.

339 learned in our treatment of binders; after all, definitions are “just” top-level binding forms!
 340 This perspective lets us simplify how we view some more baroque top-level bindings. As a
 341 contrived example, consider the following signature for a pair of top-level Agda definitions.

Listing 6 A complicated Agda signature.

```
342 private instance abstract @irr @mixed foo bar : Nat → _
```

345 In our language of binders, this definition consists of a 2-ary annotated binding of the names
 346 `foo`, `bar` that has had a sequence of binding modifiers applied to it.

347 Unfortunately, this insight does not offer a complete solution. Notably, our four target
 348 grammar differ significantly in how their treatment of type signatures. prioritize dependent
 349 pattern matching (EG: Agda, Idris 2) typically opt to have standalone type signatures: this
 350 allow for top-level pattern matches, which in turn makes it much easier to infer motives[7].
 351 Conversely, languages oriented around tactics (EG: Lean 4, Rocq) typically opt for in-line
 352 type signatures and pattern-matching expressions. This appears to be largely independent of
 353 Miranda/ML syntax split: Lean 4 borrows large parts of its syntax from Haskell, yet still
 354 opts for in-line signatures.

355 This presents us with a design decision: should our idealized grammar use inline or
 356 standalone signatures? As long as we can (easily) translate from one style to the other, we
 357 have a genuine decision to make. We have opted for the former as standalone signatures
 358 offer variations that languages with inline signatures cannot handle. As a concrete example,
 359 consider the following Agda declaration:

Listing 7 A definition with mismatched names.

```
360 id : (A : Type) → A → A
361   id B x = x
```

364 In particular, note that we have bound first argument to a different name. Translating
 365 this to a corresponding Rocq declaration then forces us to choose to use either the name from
 366 the signature or the term. Conversely, using in-line signatures does not lead us to having to
 367 make an unforced choice when translating to a separate signature, as we can simply duplicate
 368 the name in both the signature and term.

369 However, inline signatures are not completely without fault, and cause some edge cases
 370 with binding modifiers. As an example, consider the following two variants of the identity
 371 function in Agda.

Listing 8 Two variants of the identity function.

```
372 id : {A : Type} → A → A
373   id x = x
374
375 id' : {A : Type} → A → A
376   id' {A} x = x
```

379 Both definitions mark the `A` argument as an implicit, but the second definition *also* binds
 380 it in the declaration. When we pass to inline type signatures, we lose this extra layer of
 381 distinction. To account for this, we were forced to refine the visibility modifier system to
 382 distinguish between “bound” and “unbound” modifiers. This extra distinction has not proved
 383 to be too onerous in practice, and we still believe that inline signatures are the correct choice
 384 for our application.

385 We have encoded this decision in our idealized grammar by introducing a notion of a
 386 “left-hand-side” of a definition, which consists of a collection of names to be defined, and a

Reed: link to
the previous
listing

scope to define them under. This means that we view definitions like ??? not as functions $\text{id} : (A : \text{Type}) \rightarrow A \rightarrow A$ but rather as *bindings* $A : \text{Type}, x : A \vdash \text{id} : A$ in non-empty contexts. This shift in perspective has the added benefit of making the interface to other forms of parameterised definitions entirely uniform; for instance, a parameterised record is simply just a record with a non-empty left-hand side.

In Panbench, definitions and their corresponding left-hand sides are encoded via the following set of typeclasses.

Listing 9 Definitions and left-hand sides.

```

394 class Definition lhs tm defn | defn → lhs tm where
395   (.=) :: lhs → tm → defn
396
397 class DataDefinition lhs ctor defn | defn → lhs ctor where
398   data_ :: lhs → [ctor] → defn
399
400 class RecordDefinition lhs nm fld defn | defn → lhs nm fld where
401   record_ :: lhs → name → [fld] → defn
402
403

```

7 Conclusion

References

- 1 February 2026. URL: <https://github.com/GrammaticalFramework/informath>.
- 2 Ali Assaf, Guillaume Burel, Raphaël Cauderlier, David Delahaye, Gilles Dowek, Catherine Dubois, Frédéric Gilbert, Pierre Halmagrand, Olivier Hermant, and Ronan Saillard. Dedukti: a logical framework based on the $\lambda\pi$ -calculus modulo theory. (arXiv:2311.07185), November 2023. arXiv:2311.07185 [cs]. URL: <http://arxiv.org/abs/2311.07185>, doi:10.48550/arXiv.2311.07185.
- 3 Jacques Carette, Oleg Kiselyov, and Chung-Chieh Shan. Finally tagless, partially evaluated: Tagless staged interpreters for simpler typed languages. *Journal of Functional Programming*, 19(5):509–543, 2009. doi:10.1017/S0956796809007205.
- 4 Eelco Dolstra and The Nix contributors. Nix. URL: <https://github.com/NixOS/nix>.
- 5 Robert Harper, Furio Honsell, and Gordon Plotkin. A framework for defining logics. *Journal of the ACM*, 40(1):143–184, January 1993. doi:10.1145/138027.138060.
- 6 Mark P. Jones. Type classes with functional dependencies. In Gert Smolka, editor, *Programming Languages and Systems*, page 230–244, Berlin, Heidelberg, 2000. Springer. doi:10.1007/3-540-46425-5_15.
- 7 Conor McBride and James McKinna. The view from the left. *Journal of Functional Programming*, 14(1):69–111, January 2004. doi:10.1017/S0956796803004829.
- 8 Neil Mitchell. Shake before building: replacing make with haskell. *SIGPLAN Not.*, 47(9):55–66, 2012. doi:10.1145/2398856.2364538.
- 9 Aarne Ranta. *Grammatical framework: programming with multilingual grammars*. CSLI studies in computational linguistics. CSLI Publications, Center for the Study of Language and Information, Stanford, Calif, 2011.