

# Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

3 **Anonymous author**

4 Anonymous affiliation

5 **Anonymous author**

6 Anonymous affiliation

---

## 7 — Abstract —

8 We benchmark four proof assistants (Agda, Idris 2, Lean 4 and Rocq) through a single test suite.  
9 We focus our benchmarks on the basic features that all systems based on a similar foundations  
10 (dependent type theory) have in common. We do this by creating an “over language” in which to  
11 express all the information we need to be able to output *correct and idiomatic syntax* for each of our  
12 targets. Our benchmarks further focus on “basic engineering” of these systems: how do they handle  
13 long identifiers, long lines, large records, large data declarations, and so on.

14 Our benchmarks reveals both flaws and successes in all systems. We give a thorough analysis of  
15 the results.

16 We also detail the design of our extensible system. It is designed so that additional tests and  
17 additional system versions can easily be added. A side effect of this work is a better understanding  
18 of the common abstract syntactic structures of all four systems.

19 **2012 ACM Subject Classification** Replace ccsdesc macro with valid one

20 **Keywords and phrases** Add keywords

21 **Digital Object Identifier** 10.4230/LIPIcs.CVIT.2016.23

## 22 1 Introduction

23 Production-grade implementations of dependently typed programming languages are complica-  
24 tated pieces of software that implement many intricate and potentially expensive algorithms.  
25 As such, large amounts of engineering effort has been dedicated to optimizing these com-  
26 ponents. Unfortunately, engineering time is a finite resource, and this necessarily means  
27 that other parts of these systems get comparatively less attention. This often results in  
28 easy-to-miss performance problems: we have heard anecdotes from a proof assistant developer  
29 that a naïve  $O(n^2)$  fresh name generation algorithm used for pretty-printing resulted in 100x  
30 slowdowns in some pathological cases.

31 This suggests that a benchmarking suite that focuses on these simpler components  
32 could reveal some (comparatively) easy potential performance gains. Moreover, such a  
33 benchmarking suite would also be valuable for developers of new dependently typed languages,  
34 as it is much easier to optimize with a performance goal in mind. This is an instance of  
35 the classic  $m \times n$  language tooling problem: constructing a suite of  $m$  benchmarks for  $n$   
36 languages directly requires a quadratic amount of work up front, and adding either a new  
37 test case or a new language to the suite requires an additional linear amount of effort.

38 Like most  $m \times n$  tooling problems, the solution is to introduce a mediating tool. In our  
39 case, we ought to write all of the benchmarks in an intermediate language, and then translate  
40 that intermediate language to the target languages in question. There are existing languages  
41 like Dedukti[2] or Informath[1] that attempt to act as an intermediary between popular proof  
42 assistants, but these tools typically focus on translating the *content* of proofs, not exact  
43 syntactic structure. To fill this gap, we have created the Panbench system, which consists of:



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:10



LIPICS Leibniz International Proceedings in Informatics

Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

## 23:2 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

- 44 1. An extensible embedded Haskell DSL that encodes the concrete syntax of a typical  
45 dependently typed language.
- 46 2. A series of compilers for that DSL to Agda, Idris 2, Lean 4, and Rocq.
- 47 3. A benchmarking harness that can perform sandboxed builds of multiple revisions Agda,  
48 Idris 2, Lean 4, Rocq.
- 49 4. An incremental build system that can produce benchmarking reports as static HTML  
50 files or PGF plots<sup>1</sup>.

## 51 2 Methodology

52 This is really documenting the 'experiment'. The actual details of the thinking behind  
the design is in Section 6.

53 this itemized  
list should be  
expanded into  
actual text

- 54 └─ single language of tests  
55 └─ document the setup of tests, high level  
56 └─ document the setup of testing infrastructure, high level  
57 └─ linear / exponential scaling up of test 'size'

## 58 3 Results

## 59 4 Discussion

## 60 5 Infrastructure

61 One of the major goals of Panbench is to make performance analysis as low-cost as possible for  
62 language developers. Meeting this goal requires a large amount of supporting infrastructure:  
63 simply generating benchmarks is not very useful if you cannot run them nor analyze their  
64 outcomes. After some discussion, we concluded that any successful language benchmarking  
65 system should meet the following criteria:

- 66 1. It must provide infrastructure for performing sandboxed builds of compilers from source.  
67 Asking potential users to set up four different toolchains presents an extremely large  
68 barrier to adoption. Moreover, if we rely on user-provided binaries, then we have no hope  
69 of obtaining reproducible results, which in turn makes any insights far less actionable.
- 70 2. It must allow for multiple revisions of the same tool to be installed simultaneously. This  
71 enables developers to easily look for performance regressions, and quantify the impact of  
72 optimizations.
- 73 3. It must allow for multiple copies of the *same* version tool to be installed with different  
74 build configurations. This allows developers to look for performance regressions induced  
75 by different compiler versions/optimizations.
- 76 4. It must be able to be run locally on a developer's machine. Cloud-based tools are often  
77 cumbersome to use and debug, which in turn lowers adoption.
- 78 5. It must present a declarative interface for creating benchmarking environments and  
79 running benchmarks. Sandboxed builds of tools are somewhat moot if we cannot trust  
80 that a benchmark was run with the correct configuration.

---

<sup>1</sup> All plots in this paper were produced directly by Panbench!

81    6. It must present performance results in a self-contained format that is easy to understand  
82    and share. Performance statistics that require large amounts of post-processing or  
83    dedicated tools to view can not be easily shared with developers, which in turn makes  
84    the data less actionable.

85    Of these criteria, the first four present the largest engineering challenge, and are tantamount to constructing a meta-build system that is able to orchestrate *other* build systems.  
86    We approached the problem by constructing a bespoke content-addressed system atop of  
87    Shake [8], which we discuss in section 5.1. The final two criteria also presented some unforeseen  
88    difficulties, which we detail in 5.2.

## 90    5.1 The Panbench Build System

91    As noted earlier, we strongly believe that any benchmarking system should provide infrastructure  
92    for installing multiple versions of reproducibly built software. Initially, we intended  
93    to build this infrastructure for Panbench atop of Nix [4]. This is seemingly a perfect fit;  
94    after all, Nix was designed to facilitate almost exactly this use-case. However, after further  
95    deliberation, we came to the conclusion that Nix did not quite meet our needs for the  
96    following reasons:

- 97    1. Nix does not work natively on Windows. Performance problems can be operating system  
98    specific, so ruling out an OS that has a large user base that is often overlooked in testing  
99    seems unwise<sup>2</sup>.
- 100    2. Nix adds a barrier to adoption. Installing Nix is a somewhat invasive process, especially  
101    on MacOS<sup>3</sup>. We believe that it is somewhat unreasonable to ask developers to add users  
102    and modify their root directory to run a benchmarking tool, and strongly suspect that  
103    this would hamper adoption.

104    With the obvious option exhausted, we opted to create our own Nix-inspired build system  
105    based atop Shake [8]. This avoids the aforementioned problems with Nix: Shake works on  
106    Windows, and only requires potential users to install a Haskell toolchain.

107    The details of content-addressed build systems are a deep topic unto themselves, so we  
108    will only describe the key points. Systems like Nix use an *input-addressing* scheme, wherein  
109    the results of a build are stored on disk prefixed by a hash of all build inputs. Crucially,  
110    this lets the build system know where to store the result of the build before the build is run,  
111    which avoids vicious cycles where the result of a build depends on its own hash. However,  
112    most input-addressed systems also require that the hash of the inputs *solely* determines the  
113    output. On its face, this is a reasonable ask, but taking this idea to its logical conclusion  
114    requires one to remove *any* dependency on the outer environment, which in turn forces one  
115    to re-package the entirety of the software stack all the way down to `libc`. This is an admirable  
116    goal in its own right, but is actually somewhat counterproductive for our use case: there  
117    is a very real chance that we might end up benchmarking our sub-par repackaging of some  
118    obscure C dependency four layers down the stack.

119    To avoid this cascading series of dependency issues, Panbench takes a hybrid approach,  
120    wherein builds are input-addressed, but are allowed to also depend on the external environment.  
121    Additionally, the results of builds are also hashed, and stored out-of-band inside of a

---

<sup>2</sup> Currently, Panbench does not support Windows, but this is an artifact of prioritization, and not a fundamental restriction.

<sup>3</sup> The situation is even worse on x86-64 Macs, which most Nix installers simply do not support.

## 23:4 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

122 Shake database. This hash is used to perform invalidate downstream results, and also as a  
123 fingerprint to identify if two benchmarks were created from the same binary. This enables  
124 a pay-as-you go approach to reproducibility, which we hope will result in a lower adoption  
125 cost. Moreover, we can achieve fully reproducible builds by using Nix as a meta-meta build  
126 system to compile Panbench: this is how we obtained the results presented in 3.

### 127 5.2 Running Benchmarks and Generating Reports

## 128 6 The Design of Panbench

129 At its core, Panbench is a tool for producing grammatically well-formed concrete syntax  
130 across multiple different languages. Crucially, Panbench does *not* require that the syntax  
131 produced is well-typed or even well-scoped: if it did, then it would be impossible to benchmark  
132 how systems perform when they encounter errors. This seemingly places Panbench in stark  
133 contrast with other software tools for working with the meta-theoretic properties of type  
134 systems, which are typically concerned only with well-typed terms.

135 However, core task of Panbench is not that different from the task of a logical framework [5]:  
136 Both systems exist to manipulate judgements, inference rules, and derivations: Panbench  
137 just works with *grammatical* judgements and production rules rather than typing judgments  
138 and inference rules. In this sense Panbench is a *grammatical* framework<sup>4</sup> rather than a logical  
139 one.

140 This similarity let us build Panbench atop well-understood design principles. In particular,  
141 a mechanized logical framework typically consists of two layers:

- 142 1. A layer for defining judgements à la relations.
- 143 2. A logic programming layer for synthesizing derivations.

144 To use a logical framework, one first encodes a language by laying out all of the judgements.  
145 Then, one needs to prove an adequacy theorem on the side that shows that their encoding of  
146 the judgements actually aligns with the language. However, if one wanted to mechanize this  
147 adequacy proof, then a staged third layer that consists of a more traditional proof assistant  
148 would be required.

149 If we take this skeleton of a design and transpose it to work with grammatical constructs  
150 rather than logical ones, we will also obtain three layers:

- 151 1. A layer for defining grammars à la relations.
- 152 2. A logic programming layer for synthesizing derivations.
- 153 3. A staged functional programming layer for proving “adequacy” results.

154 In this case, an adequacy result for a given language  $\mathcal{L}$  is a constructive proof that all  
155 grammatical derivations written within the framework can be expressed within the concrete  
156 syntax of a language  $\mathcal{L}$ . However, the computational content of such a proof essentially  
157 amounts to a compiler written in the functional programming layer.

---

<sup>4</sup> Not to be confused with *the Grammatical Framework* [9], which aims to be a more natural-language focused meta-framework for implementing and translating between grammars.

158 **6.1 Implementing The Grammatical Framework**

159 Implementing a bespoke hybrid of a logic and functional programming language is no small  
 160 feat, and also requires prospective users to learn yet another single-purpose tool. Luckily,  
 161 there already exists a popular, industrial-grade hybrid logic/functional programming language  
 162 in wide use: GHC Haskell.

163 At first glance, Haskell does not contain a logic programming language. However, if we  
 164 enable enough GHC extensions, the typeclass system can be made *just* rich enough to encode  
 165 a simply-typed logical framework. The key insight is that we can encode each production  
 166 rule using multi-parameter type classes with a single method. Moreover, we can encode our  
 167 constructive adequacy proofs for a given set of production rules as instances that translate  
 168 each of the productions in the abstract grammar to productions in the syntax of an actual  
 169 language.

170 As a concrete example, consider the grammar of the following simple imperative language.

```
172 <expr> := x
173   | n
174   | <expr> '+' <expr>
175   | <expr> '*' <expr>
176 <stmt> := <var> '=' <expr>
177   | 'while' <expr> 'do' <stmt>
178   | <stmt> ';' <stmt>
```

180 We can then encode this grammar with the following set of multi-parameter typeclasses:

**Listing 1** An example tagless encoding.

```
181 class Var expr where
182   var :: String → expr
183
184 class Lit expr where
185   lit :: Int → expr
186
187 class Add expr where
188   add :: expr → expr → expr
189
190 class Mul expr where
191   mul :: expr → expr → expr
192
193 class Assign expr stmt where
194   assign :: String → expr → stmt
195
196 class While expr stmt where
197   while :: expr → stmt → stmt
198
199 class AndThen stmt where
200   andThen :: stmt → stmt → stmt
```

203 This style of language encoding is typically known as the untyped variant of *finally*  
 204 *tagless*[3], and is well-known technique. However, our encoding is a slight refinement of  
 205 the usual tagless style. In particular, we restrict ourselves to a single class per production  
 206 rule, whereas other tagless encodings often use a class per syntactic category. This more  
 207 fine-grained approach allows us to encode grammatical constructs that are only supported  
 208 by a subset of our target grammars; see section 6.2 for examples.

## 23:6 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

209        Unfortunately, the encoding above has some serious ergonomic issues. In particular,  
210      expressions like `assign "x"` (`lit 4`) will result in an unsolved metavariable for `expr`, as there  
211      may be multiple choices of `expr` to use when solving the `Assign ?expr stmt` constraint. Luckily,  
212      we can resolve ambiguities of this form through judicious use of functional dependencies[6],  
213      as demonstrated below.

214      **Listing 2** A tagless encoding with functional dependencies.  
215      

```
class Assign expr stmt | stmt → expr where
216        assign :: String → expr → stmt
217
218      class While expr stmt | stmt → expr where
219        while :: expr → stmt → stmt
```

  
220

## 221 6.2 Designing The Language

222      Now that we've fleshed out how we are going to encode our grammatical framework into  
223      our host language, it's time to design our idealized abstract grammar. All of our target  
224      languages roughly agree on a subset of the grammar of non-binding terms: the main sources  
225      of divergence are binding forms and top-level definitions<sup>5</sup>. This is ultimately unsurprisingly:  
226      dependent type theories are fundamentally theories of binding and substitution, so we would  
227      expect some variation in how our target languages present the core of their underlying  
228      theories.

229      This presents an interesting language design problem. Our idealized grammar will need to  
230      find some syntactic overlap between all of our four target languages. Additionally, we would  
231      also like for our solution to be (reasonably) extensible. Finding the core set of grammatical  
232      primitives to accomplish this task is surprisingly tricky, and requires a close analysis of fine  
233      structure of binding.

### 234 6.2.1 Binding Forms

235      As users of dependently typed languages are well aware, a binding form carries much more  
236      information than just a variable name and a type. Moreover, this extra information can have  
237      a large impact on typechecking performance, as is the case with implicit/visible arguments.  
238      To make matters worse, languages often offer multiple syntactic options for writing the same  
239      binding form, as is evidenced by the prevalence of multi-binders like  $(x\ y\ z : A) \rightarrow B$ . Though  
240      such binding forms are often equivalent to their single-binder counterparts as *abstract* syntax,  
241      they may have different performance characteristics, so we cannot simply lower them to a  
242      uniform single-binding representation. To account for these variations, we have designed a  
243      sub-language dedicated solely to binding forms. This language classifies the various binding  
244      features along three separate axes: binding arity, binding annotations, and binding modifiers.

245      Binding arities and annotations are relatively self-explanatory, and classify the number of  
246      names bound, along with the type of annotation allowed. Our target languages all have their  
247      binding arities falling into one of three classes: *n*-ary, unary, or nullary. We can similarly  
248      characterise annotations into three categories: required, optional, or forbidden.

---

<sup>5</sup> As we shall see in section 6.2.2, the syntactical divergence of top-level definitions is essentially about binders as well.

249 This language of binders is encoded in the implementation as a single class `Binder` that is  
 250 parameterised by higher-kinded types `arity :: Type -> Type` and `ann :: Type -> Type`, and  
 251 provide standardized types for all three binding arities and annotations.

**Listing 3** The basic binding constructs in Panbench.

```
252 class Binder arity nm ann tm cell | cell -> nm tm where
253   binder :: arity nm -> ann tm -> cell
254
255   -- | No annotation or arity.
256   data None nm = None
257
258   -- | A single annotation or singular arity.
259   newtype Single a = Single { unSingle :: a }
260
261   -- | Multi-binders.
262   type Multi = []
263
264   -- | Infix operator for an annotated binder with a single name.
265   (. :) :: (Binder Single nm Single tm cell) => nm -> tm -> cell
266   nm .: tp = binder (Single nm) (Single tp)
267
268   -- | Infix operator for an annotated binder.
269   (. :*) :: (Binder arity nm Single tm cell) => arity nm -> tm -> cell
270   nms .: tp = binder nms (Single tp)
271
```

273 Production rules that involve binding forms are encoded as classes that are parametric over  
 274 a notion of a binding cell, as demonstrated below.

**Listing 4** The Panbench class for II-types.

```
275 class Pi cell tm | tm -> cell where
276   pi :: [cell] -> tm -> tm
277
```

279 Decoupling the grammar of binding forms from the grammar of binders themselves allows  
 280 us to be somewhat polymorphic over the language of binding forms when writing generators.  
 281 This in turn means that we can potentially re-use generators when extending Panbench with  
 282 new target grammars that may support only a subset of the binding features present in our  
 283 four target grammars.

284 Binding modifiers, on the other hand, require a bit more explanation. A binding modifier  
 285 captures features like implicit arguments, which do not change the number of names bound  
 286 nor their annotations, but rather how those bound names get treated by the rest of the  
 287 system. Currently, Panbench only supports visibility-related modifiers, but we have designed  
 288 the system so that it is easy to extend with new modifiers; EG: quantities in Idris 2 or  
 289 irrelevance annotations in Agda.

290 The language of binding modifiers is implemented as the following set of Haskell type-  
 291 classes.

**Listing 5** Typeclasses for binding modifiers.

```
292 class Implicit cell where
293   implicit :: cell -> cell
294
295 class SemiImplicit cell where
296   semiImplicit :: cell -> cell
297
```

299 This is a case where the granular tagless encoding shines. Both Lean 4 and Rocq have a  
 300 form of semi-implicits<sup>6</sup>, whereas Idris 2 and Agda have no such notion. Decomposing the  
 301 language of binding modifiers into granular pieces lets us write benchmarks that explicitly  
 302 require support for features like semi-implicits. Had we used a monolithic class that encodes  
 303 the entire language of modifiers, we would have to resort to runtime errors (or, even worse,  
 304 dubious attempts at translation).

### 305 6.2.2 Top-Level Definitions

306 The question of top-level definitions is much thornier, and there seems to be less agreement  
 307 on how they ought to be structured. Luckily, we can re-apply many of the lessons we  
 308 learned in our treatment of binders; after all, definitions are “just” top-level binding forms!  
 309 This perspective lets us simplify how we view some more baroque top-level bindings. As a  
 310 contrived example, consider the following signature for a pair of top-level Agda definitions.

**311 Listing 6** A complicated Agda signature.

```
312 private instance abstract @irr @mixed foo bar : Nat → _
```

314 In our language of binders, this definition consists of a 2-ary annotated binding of the names  
 315 `foo`, `bar` that has had a sequence of binding modifiers applied to it.

316 Unfortunately, this insight does not offer a complete solution. Notably, our four target  
 317 grammar differ significantly in how their treatment of type signatures. prioritize dependent  
 318 pattern matching (EG: Agda, Idris 2) typically opt to have standalone type signatures: this  
 319 allow for top-level pattern matches, which in turn makes it much easier to infer motives[7].  
 320 Conversely, languages oriented around tactics (EG: Lean 4, Rocq) typically opt for in-line  
 321 type signatures and pattern-matching expressions. This appears to be largely independent of  
 322 Miranda/ML syntax split: Lean 4 borrows large parts of its syntax from Haskell, yet still  
 323 opts for in-line signatures.

324 This presents us with a design decision: should our idealized grammar use inline or  
 325 standalone signatures? As long as we can (easily) translate from one style to the other, we  
 326 have a genuine decision to make. We have opted for the former as standalone signatures  
 327 offer variations that languages with inline signatures cannot handle. As a concrete example,  
 328 consider the following Agda declaration:

**329 Listing 7** A definition with mismatched names.

```
330 id : (A : Type) → A → A
331 id B x = x
```

333 In particular, note that we have bound first argument to a different name. Translating  
 334 this to a corresponding Rocq declaration then forces us to choose to use either the name from  
 335 the signature or the term. Conversely, using in-line signatures does not lead us to having to  
 336 make an unforced choice when translating to a separate signature, as we can simply duplicate  
 337 the name in both the signature and term.

338 However, inline signatures are not completely without fault, and cause some edge cases  
 339 with binding modifiers. As an example, consider the following two variants of the identity  
 340 function in Agda.

---

<sup>6</sup> We use the term *semi-implicit* argument to refer to an implicit argument that is not eagerly instantiated. In Rocq these are known as non-maximal implicits.

**Listing 8** Two variants of the identity function.

```

341   id : {A : Type} → A → A
342   id x = x
343
344
345   id' : {A : Type} → A → A
346   id' {A} x = x

```

Both definitions mark the `A` argument as an implicit, but the second definition *also* binds it in the declaration. When we pass to inline type signatures, we lose this extra layer of distinction. To account for this, we were forced to refine the visibility modifier system to distinguish between “bound” and “unbound” modifiers. This extra distinction has not proved to be too onerous in practice, and we still believe that inline signatures are the correct choice for our application.

We have encoded this decision in our idealized grammar by introducing a notion of a “left-hand-side” of a definition, which consists of a collection of names to be defined, and a scope to define them under. This means that we view definitions like `??? : (A : Type) → A → A` but rather as *bindings* `A : Type, x : A ⊢ id : A` in non-empty contexts. This shift in perspective has the added benefit of making the interface to other forms of parameterised definitions entirely uniform; for instance, a parameterised record is simply just a record with a non-empty left-hand side.

In Panbench, definitions and their corresponding left-hand sides are encoded via the following set of typeclasses.

Reed: link to the previous listing

**Listing 9** Definitions and left-hand sides.

```

363 class Definition lhs tm defn | defn → lhs tm where
364   (.=) :: lhs → tm → defn
365
366 class DataDefinition lhs ctor defn | defn → lhs ctor where
367   data_ :: lhs → [ctor] → defn
368
369 class RecordDefinition lhs nm fld defn | defn → lhs nm fld where
370   record_ :: lhs → name → [fld] → defn
371

```

## 7 Conclusion

### References

- 1 February 2026. URL: <https://github.com/GrammaticalFramework/informatH>.
- 2 Ali Assaf, Guillaume Burel, Raphaël Cauderlier, David Delahaye, Gilles Dowek, Catherine Dubois, Frédéric Gilbert, Pierre Halmagrand, Olivier Hermant, and Ronan Saillard. Dedukti: a logical framework based on the  $\lambda\pi$ -calculus modulo theory. (arXiv:2311.07185), November 2023. arXiv:2311.07185 [cs]. URL: <http://arxiv.org/abs/2311.07185>, doi:10.48550/arXiv.2311.07185.
- 3 Jacques Carette, Oleg Kiselyov, and Chung-Chieh Shan. Finally tagless, partially evaluated: Tagless staged interpreters for simpler typed languages. *Journal of Functional Programming*, 19(5):509–543, 2009. doi:10.1017/S0956796809007205.
- 4 Eelco Dolstra and The Nix contributors. Nix. URL: <https://github.com/NixOS/nix>.
- 5 Robert Harper, Furio Honsell, and Gordon Plotkin. A framework for defining logics. *Journal of the ACM*, 40(1):143–184, January 1993. doi:10.1145/138027.138060.

## 23:10 Panbench: A Comparative Benchmarking Tool for Dependently-Typed Languages

- 387    6    Mark P. Jones. Type classes with functional dependencies. In Gert Smolka, editor, *Programming  
388                  Languages and Systems*, page 230–244, Berlin, Heidelberg, 2000. Springer. doi:10.1007/  
389                  3-540-46425-5\_15.
- 390    7    Conor McBride and James McKinna. The view from the left. *Journal of Functional Programming*,  
391                  14(1):69–111, January 2004. doi:10.1017/S0956796803004829.
- 392    8    Neil Mitchell. Shake before building: replacing make with haskell. *SIGPLAN Not.*, 47(9):55–66,  
393                  2012. doi:10.1145/2398856.2364538.
- 394    9    Aarne Ranta. *Grammatical framework: programming with multilingual grammars*. CSLI  
395                  studies in computational linguistics. CSLI Publications, Center for the Study of Language and  
396                  Information, Stanford, Calif, 2011.